

HW4

March 2, 2015

1 STAT 541 HW4

The following code is Python. I checked it against the SAS output. Seems legit.

1.1 Problem 1

```
In [52]: import pandas as pd
import numpy as np
import statsmodels.api as sm
from statsmodels.formula.api import ols

### enter data
prob1_data = pd.DataFrame({'temperature' : np.repeat(['1520', '1600', '1680', '1760',\
'1840'], 6),
                           'fail_time' : [1253, 1435, 1771, 4027, 5434, 5614,
1190, 1286, 1550, 2125, 2557, 2845,
751, 837, 848, 1038, 1361, 1443,
611, 691, 751, 772, 808, 859,
513, 546, 517, 420, 471, 556]})

### find logged group means and std
logmeans = np.log(prob1_data.groupby('temperature').mean())
logstd = np.log(prob1_data.groupby('temperature').std())

### fit model
logmeans = sm.add_constant(logmeans)
prob1_model = sm.OLS(logstd, logmeans).fit()

### calculate lambda
lambdaa = 1 - prob1_model.params[1]
lambdaa
```

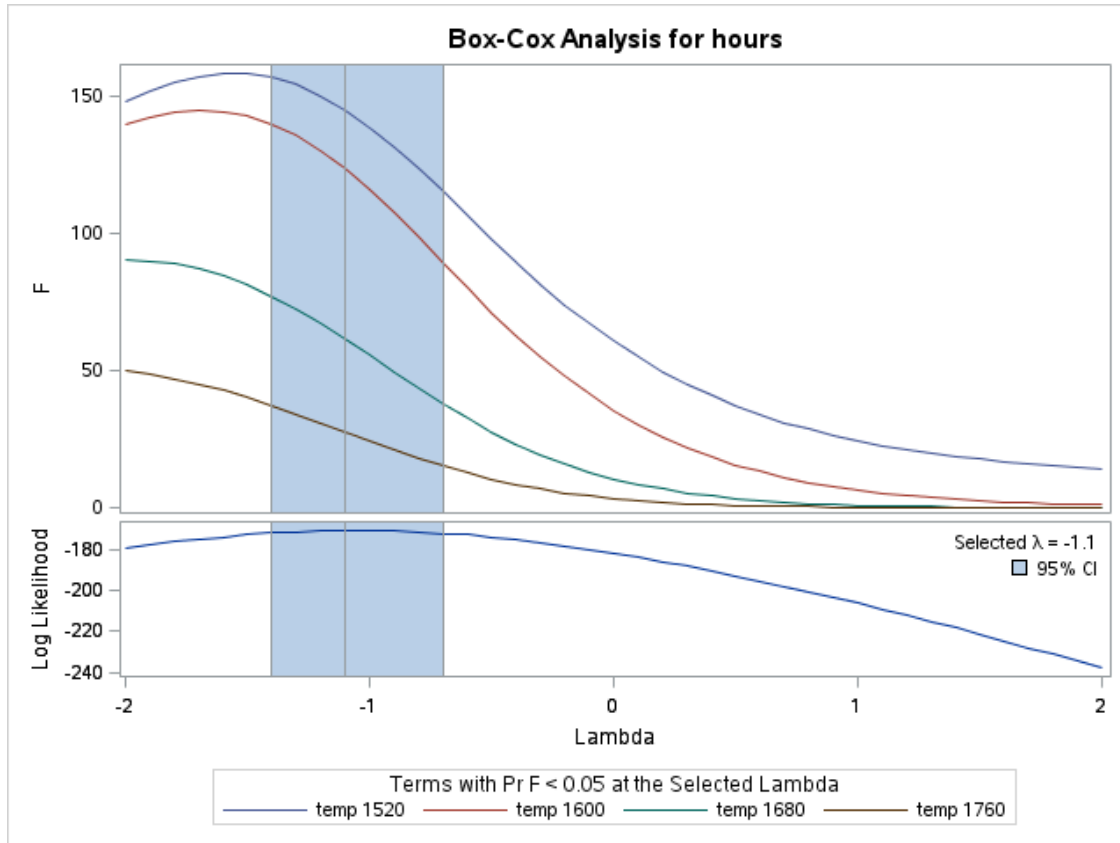
```
Out[52]: -0.99756363490905109
```

The empirical method results in $\lambda \approx -1$ which is approximately a reciprocal transformation.

1.2 Problem 2

```
In [53]: from IPython.display import Image
Image(filename='hw4_boxcox.png')
```

```
Out[53]:
```



The Box-Cox function in Python doesn't work for the purposes of this context. The SAS output says that the recommended transformation is roughly the reciprocal transformation as well ($\lambda \approx -1.1$).

1.3 Problem 3

1.3.1 Part (a)

The means model can be notated as

$$y_{ij} = \mu_i + \epsilon_{ij} \quad 1 \leq i \leq 5, 1 \leq j \leq 6$$

$$\epsilon_{ij} \stackrel{iid}{\sim} N(0, \sigma^2)$$

where i is the index for group and j is the index for observations within groups. The parameter μ_i represents the mean failure time in hours of group i , ϵ_{ij} is the error term for observation j in group i , and σ^2 is the variance of the errors.

1.3.2 Part (b)

In [69]: `from scipy import stats`

```
### seperate data by temperature level
g1=prob1_data.query('temperature == "1520"').fail_time
g2=prob1_data.query('temperature == "1600"').fail_time
g3=prob1_data.query('temperature == "1680"').fail_time
g4=prob1_data.query('temperature == "1760"').fail_time
```

```
g5=prob1_data.query('temperature =="1840"').fail_time
```

```
### run levene's test
```

```
levене_out = stats.levene(g1, g2, g3, g4, g5, center='mean')
{'F-statistic' : levene_out[0], 'p-value' : levene_out[1]}
```

```
Out[69]: {'F-statistic': 37.404536440145407, 'p-value': 3.2862755117618303e-10}
```

The results of Levene's test will inform a hypothesis test of

$$H_0 : \sigma_i = \sigma_j \quad \text{for all } i, j$$

$$H_a : \sigma_i \neq \sigma_j \quad \text{for some } i, j$$

but Levene's test is actually testing the following hypotheses: let μ_i^* denote the mean $\hat{\mu}_i^*$ absolute deviation for group i which is estimated by $\hat{\mu}_i^* = \frac{1}{n} \sum_j |y_{ij} - \bar{y}_i|$

$$H_0 : \mu_i^* = \mu_j^* \quad \text{for all } i, j$$

$$H_a : \mu_i^* \neq \mu_j^* \quad \text{for some } i, j$$

The F-statistic of about 37.4 results in a p-value of .000.... This provides plenty of evidence against the null hypothesis and we conclude that the variances in failure time between the different temperature levels are not equal.

1.3.3 Part (c)


```
In [75]: import matplotlib.pyplot as plt
         %matplotlib inline
```

```
### fit model and get residuals and fitted values
```

```
prob1_model3 = ols('fail_time ~ temperature', prob1_data).fit()
resid = prob1_model3.resid
fitted = prob1_model3.fittedvalues
```

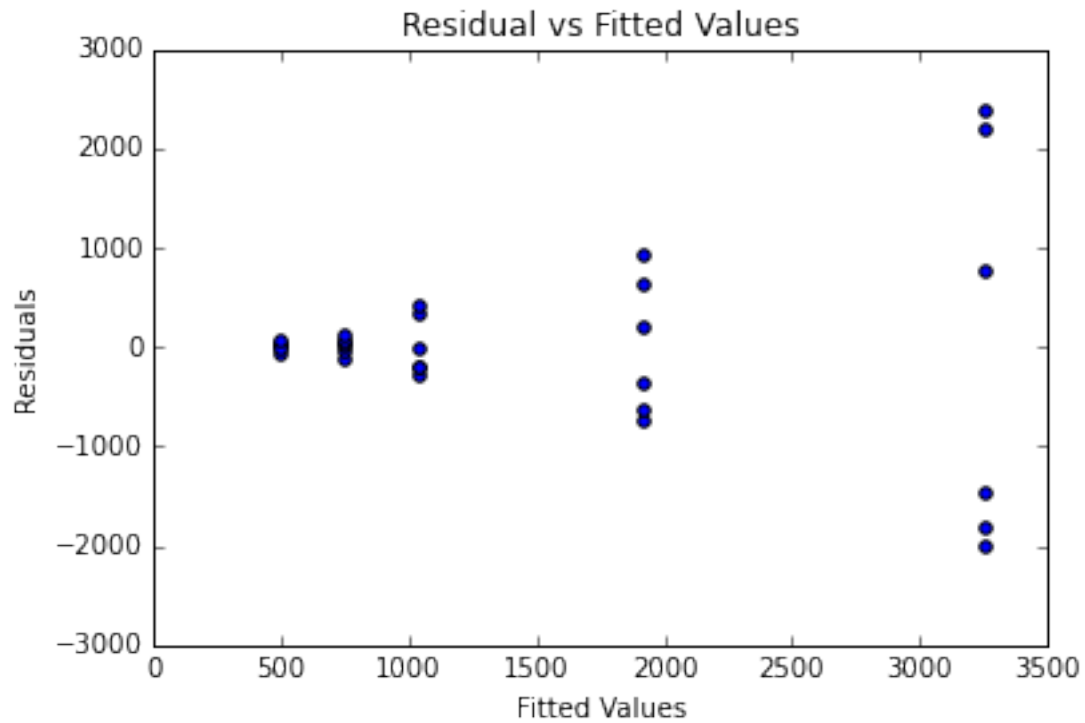
```
### normal probability plot
```

```
sm.qqplot(resid, line='s')
title('Q-Q Plot')
plt.show()
```



In the normal probability plot the points don't lie closely to the diagonal line indicating that the residuals are not normally distributed.

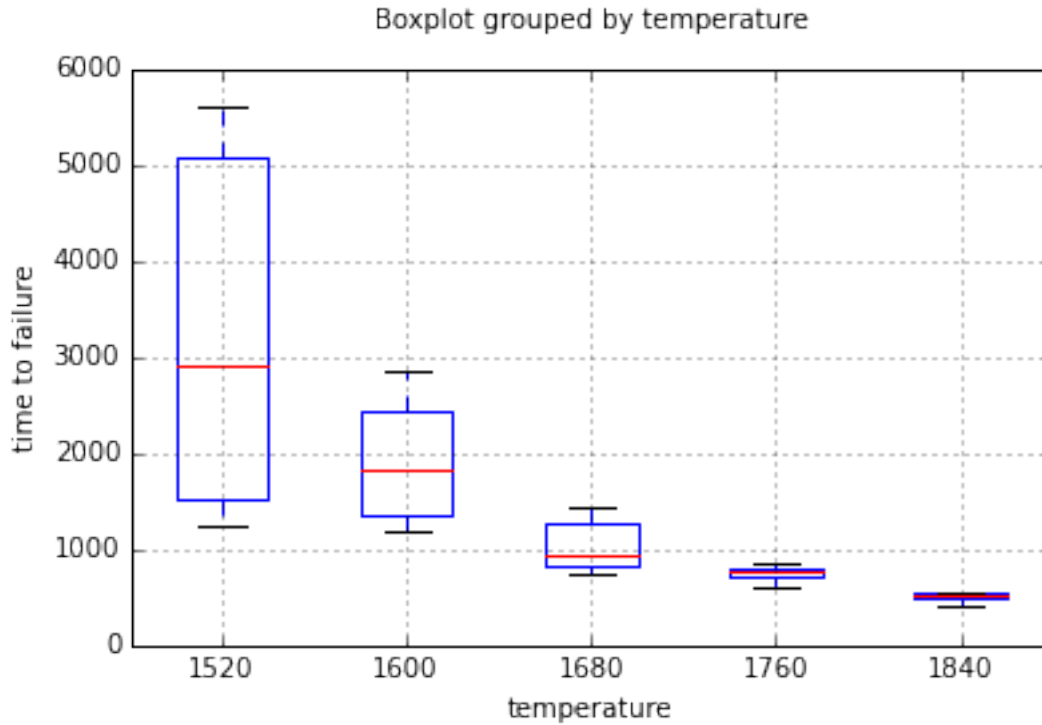
```
In [57]: ### fitted values vs residuals plot
plt.scatter(fitted, resid)
xlabel('Fitted Values')
ylabel('Residuals')
title('Residual vs Fitted Values')
plt.show()
```



In the residuals versus predicted value plot there is an obvious fanning of variance as the fitted values increase. This implies that the homogeneity of variance assumption is violated.

1.3.4 Part (d)

```
In [61]: prob1_data.boxplot(column = 'fail_time', by = 'temperature')
        ylabel('time to failure')
        title('')
        plt.show()
```



It seems pretty clear that as the mean failure time increases, so does the variance for failure time per temperature grouping.

1.3.5 Pard (e)

```
In [63]: from statsmodels.stats.multicomp import MultiComparison
```

```
tukey_test = MultiComparison(probl_data['fail_time'], probl_data['temperature'])
print tukey_test.tukeyhsd()
```

Multiple Comparison of Means - Tukey HSD,FWER=0.05

group1	group2	meandiff	lower	upper	reject
0	1	-1330.1667	-2966.6521	306.3187	False
0	2	-2209.3333	-3845.8187	-572.8479	True
0	3	-2507.0	-4143.4854	-870.5146	True
0	4	-2751.8333	-4388.3187	-1115.3479	True
1	2	-879.1667	-2515.6521	757.3187	False
1	3	-1176.8333	-2813.3187	459.6521	False
1	4	-1421.6667	-3058.1521	214.8187	False
2	3	-297.6667	-1934.1521	1338.8187	False
2	4	-542.5	-2178.9854	1093.9854	False
3	4	-244.8333	-1881.3187	1391.6521	False

It appears that the 1520 temperature group is significantly different from the 1680, 1760, and 1840 groups and all other group means are not considered significantly different.

1.4 Problem 4

1.4.1 Part (a)

```
In [64]: prob1_data['fail_time_recip'] = 1/prob1_data['fail_time']
```

The model is similar which only the response being transformed. Let $w_{ij} = \frac{1}{y_{ij}}$. The means model can be notated as

$$w_{ij} = \mu_i + \epsilon_{ij} \quad 1 \leq i \leq 5, 1 \leq j \leq 6$$

$$\epsilon_{ij} \stackrel{iid}{\sim} N(0, \sigma^2)$$

where i is the index for group and j is the index for observations within groups. The parameter μ_i represents the mean reciprical failure time in hours of group i , ϵ_{ij} is the error term for observation j in group i , and σ^2 is the variance of the errors.

1.4.2 Part (b)

```
In [68]: ### seperate data by temperature level
g1_recip=prob1_data.query('temperature == "1520"').fail_time_recip
g2_recip=prob1_data.query('temperature == "1600"').fail_time_recip
g3_recip=prob1_data.query('temperature == "1680"').fail_time_recip
g4_recip=prob1_data.query('temperature == "1760"').fail_time_recip
g5_recip=prob1_data.query('temperature == "1840"').fail_time_recip

### run levene's test
levене_out_recip = stats.levene(g1_recip, g2_recip, g3_recip, g4_recip, g5_recip, center='mean'
{'F-statistic' : levene_out_recip[0], 'p-value' : levene_out_recip[1]})
```

```
Out[68]: {'F-statistic': 1.3373187128746151, 'p-value': 0.28378015800416206}
```

The results of Levene's test will inform a hypothesis test of

$$H_0 : \sigma_i = \sigma_j \quad \text{for all } i, j$$

$$H_a : \sigma_i \neq \sigma_j \quad \text{for some } i, j$$

but Levene's test is actually testing the following hypotheses: let μ_i^{**} denote the mean absolute deviation for group i (estimated by $\mu_i^{**} = \frac{1}{n} \sum_j |w_{ij} - \bar{w}_i|$)

$$H_0 : \mu_i^{**} = \mu_j^{**} \quad \text{for all } i, j$$


$$H_a : \mu_i^{**} \neq \mu_j^{**} \quad \text{for some } i, j$$

The F-statistic of about 1.337 results in a p-value of .284. This provides little evidence that after taking the reciprical of the failure times that there is a difference in variance between groups.

1.4.3 Part (c)

```
In [76]: ### fit model and get residuals and fitted values
prob1_model4 = ols('fail_time_recip ~ temperature', prob1_data).fit()
resid_recip = prob1_model4.resid
fitted_recip = prob1_model4.fittedvalues


### normal probability plot
sm.qqplot(resid_recip, line='s')
title('Q-Q Plot')
plt.show()
```



The normal probability plot looks pretty good. The dots lie close to the diagonal line, indicating that the residuals are approx normally distributed.

```
In [73]: ### residuals vs fitted value plot
plt.scatter(fitted_recip, resid_recip)

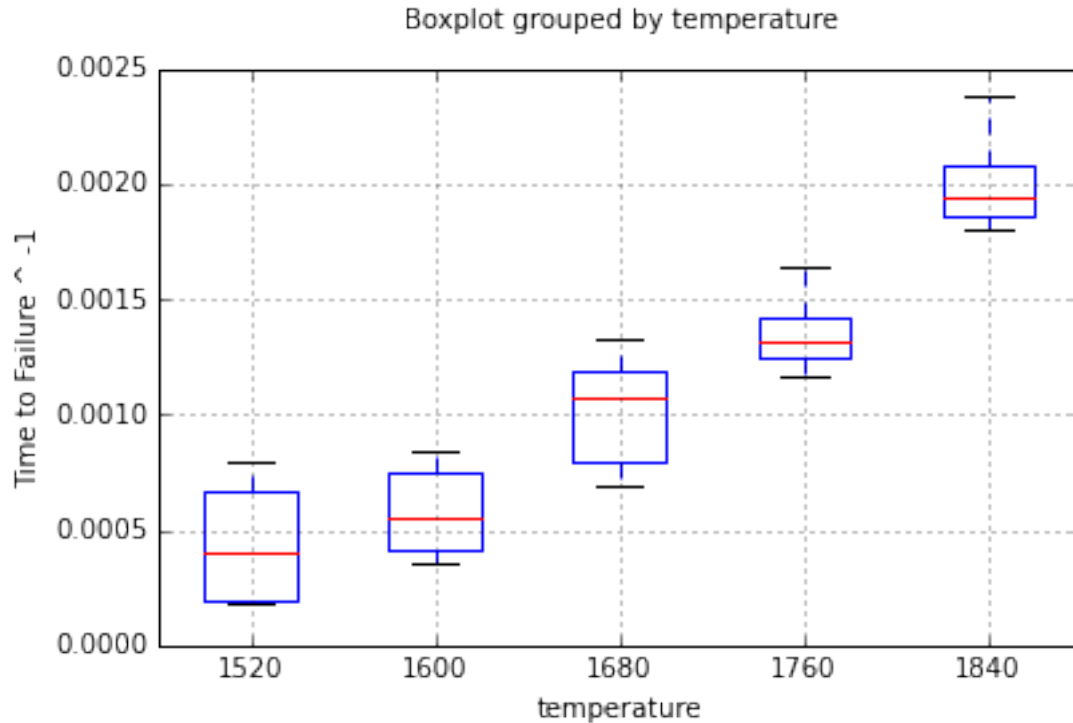
### because range of values is so small, the plot has to be adjusted
range_x = max(fitted_recip) - min(fitted_recip)
range_y = max(resid_recip) - min(resid_recip)
plt.xlim(min(fitted_recip) - .1*range_x, max(fitted_recip) + .1*range_x)
plt.ylim(min(resid_recip) - .1*range_y, max(resid_recip) + .1*range_y)
xlabel('Fitted Values')
ylabel('Residuals')
title('Residuals vs Fitted Values')
plt.show()
```

In the residual versus fitted value plot there is no obvious fanning implying that there is no homogeneity of variance violation.

1.4.4 Part (d)

```
In [78]: prob1_data.boxplot(column = 'fail_time_recip', by = 'temperature')
        ylabel('Time to Failure ^ -1')
        title('')
        plt.show()
```



It appears that the variance remains constant across the different temperature groups, regardless of mean reciprocal time to failure.

1.4.5 Part (e)

```
In [79]: tukey_test_recip = MultiComparison(probl_data['fail_time_recip'], probl_data['temperature'])
         print tukey_test_recip.tukeyhsd()
```

Multiple Comparison of Means - Tukey HSD,FWER=0.05

group1	group2	meandiff	lower	upper	reject
0	1	0.0001	-0.0003	0.0005	False
0	2	0.0006	0.0002	0.001	True
0	3	0.0009	0.0005	0.0013	True
0	4	0.0016	0.0012	0.0019	True
1	2	0.0004	0.0	0.0008	True
1	3	0.0008	0.0004	0.0012	True
1	4	0.0014	0.001	0.0018	True
2	3	0.0003	-0.0001	0.0007	False
2	4	0.001	0.0006	0.0014	True
3	4	0.0007	0.0003	0.001	True

It appears that the 1520 temperature group is not significantly different from the 1600 group and the 1680 group is not significantly different from the 1760 group. All other group comparisons are considered sig different.

1.5 Problem 5

I would definitely go with the transformed data. The diagnostic plots look much better indicating that we can assume the model assumptions are met. Looking at the original boxplot of the untransformed data, it appears that there were actually more differences than what the Tukey comparisons suggested. By taking the reciprocal of the failure times we have ‘tamed’ the increasing nature of spread which allows us to see more accurately differences in our analysis.

1.6 Problem 6

```
In [82]: prob1_model5 = ols('fail_time_recip ~ C(temperature, Poly)', prob1_data).fit()
prob1_model5.summary()
```

```
Out[82]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                        OLS Regression Results
=====
Dep. Variable:          fail_time_recip      R-squared:                0.879
Model:                  OLS                 Adj. R-squared:            0.859
Method:                 Least Squares        F-statistic:              45.20
Date:                   Sun, 01 Mar 2015      Prob (F-statistic):       4.31e-11
Time:                   23:00:30              Log-Likelihood:           211.59
No. Observations:       30                   AIC:                     -413.2
Df Residuals:           25                   BIC:                     -406.2
Df Model:                4
=====
                        coef      std err      t      P>|t|      [95.0% Conf. Int.]
-----
Intercept                0.0011    4.19e-05    25.782    0.000         0.001
C(temperature, Poly).Linear  0.0012    9.36e-05    13.139    0.000         0.001
C(temperature, Poly).Quadratic  0.0002    9.36e-05     2.662    0.013        5.64e-05
C(temperature, Poly).Cubic   3.962e-06  9.36e-05     0.042    0.967        -0.000
C(temperature, Poly)^4       9.787e-05  9.36e-05     1.046    0.306       -9.49e-05
=====
Omnibus:                 6.560    Durbin-Watson:             1.315
Prob(Omnibus):            0.038    Jarque-Bera (JB):          2.202
Skew:                     0.240    Prob(JB):                  0.332
Kurtosis:                 1.762    Cond. No.:                 2.24
=====
"""
```

As expected, it looks like there is a linear trend ($t = 13.139$, $p\text{-value} = .001$) and there is evidence for curvature ($t = 2.662$, $p\text{-value} = .013$). There is not any evidence for a cubic or quartic term.

1.7 Problem 7

1.7.1 Part (a)

I don't think it is reasonable to perform a WLS ANOVA because it appears that there is a relationship between mean response and variability. If there wasn't such a relationship then WLS would be an alright choice. Using WLS would not capture this relationship.

1.7.2 Part (b)

I would say performing a Box-Cox transformation does sound reasonable as it is apparent that the variance is proportional to some power of the mean response per group. The Box-Cox procedure seeks to find this power.

1.8 Problem 8

```
In [104]: import math
import statsmodels.stats.power as smp

sigma = 25
mus = [60., 50., 60., 55.]
mu = np.mean(mus)
f = math.sqrt(((60 - mu)**2/4 + (50 - mu)**2/4 + (60 - mu)**2/4 + (55 - mu)**2/4)/sigma)

seq1 = [.9, .91, .92, .93, .94]
sample_size = [smp.FTestAnovaPower().solve_power(effect_size=f, alpha=0.05, power=x, k_groups=4)
                for x in seq1]
pd.DataFrame({'power (1 - beta)': seq1, 'Group Sample Size (n)': sample_size})
```

Out[104]:

	Group Sample Size (n)	power (1 - beta)
0	6.231466	0.90
1	6.400805	0.91
2	6.587819	0.92
3	6.797018	0.93
4	7.035120	0.94

It appears that for each of the four groups, sample sizes of $n = 7$ are required to have a power of at least $1 - \beta = .9$.

1.9 Problem 9

1.9.1 Part (a)

```
In [105]: sigma = 36
mus = [60., 50., 60., 55.]
mu = np.mean(mus)
f = math.sqrt(((60 - mu)**2/4 + (50 - mu)**2/4 + (60 - mu)**2/4 + (55 - mu)**2/\
4)/sigma)

seq1 = [.9, .91, .92, .93, .94]
sample_size2 = [smp.FTestAnovaPower().solve_power(effect_size=f, alpha=0.05, power=x, k_group=4)
                 for x in seq1]
pd.DataFrame({'power (1 - beta)': seq1, 'Group Sample Size (n)': sample_size2})
```

Out[105]:

	Group Sample Size (n)	power (1 - beta)
0	8.470832	0.90
1	8.715748	0.91
2	8.986171	0.92
3	9.288634	0.93
4	9.632826	0.94

It appears that for each of the four groups, sample sizes of $n = 9$ are required to have a power at least $1 - \beta = .9$.

1.9.2 Part (b)

```
In [106]: sigma = 49
mus = [60., 50., 60., 55.]
mu = np.mean(mus)
f = math.sqrt(((60 - mu)**2/4 + (50 - mu)**2/4 + (60 - mu)**2/4 + (55 - mu)**2/\
4)/sigma)
```

```
seq1 = [.9, .91, .92, .93, .94]
sample_size2 = [smp.FTestAnovaPower().solve_power(effect_size=f, alpha=0.05, power=x, k_group=5) for x in seq1]
pd.DataFrame({'power (1 - beta)' : seq1, 'Group Sample Size (n)' : sample_size2})
```

```
Out[106]:
```

	Group Sample Size (n)	power (1 - beta)
0	11.132370	0.90
1	11.466476	0.91
2	11.835269	0.92
3	12.247781	0.93
4	12.717156	0.94

It appears that for each of the four groups, sample sizes of $n = 12$ are required to have a power at least $1 - \beta = .9$

1.9.3 Parts (c) and (d)

As the variability increases, the sensitivity of an ANOVA F-test decreases. Thus, the need for more data to detect an actual difference is required. As σ increases, so should the sample sizes n .